الجامعة السورية الخاصة
SYRIAN PRIVATE UNIVERSITY

Syrian Private University

College Of Artificial Intelligence

AI-FireGuard

Thermal- RGB Fusion for UAV-based Wildfire Monitoring

Prepared By

Bashar Alsaleh

Khawla Al-Shaikh

Judy Dukmak

Supervised By

Dr. Riad Sonbol

Eng. Aya Alaswad

Contents

Abstract 6
تمهيد 7
Chapter One 8
Introduction 8
Abstract 8
1.2 Causes 9
1.3 Results 9
1.4 Fact 10
Chapter Two 11
Literature Reviews 11
2.1 Introduction 11
2.2 Dataset Classification Based on Modalities and Acquisition Conditions 11
2.2.1 RGB–Thermal UAV Datasets 11
2.2.2 RGB-Only UAV Datasets 12
2.3 Model Architectures and Fusion Strategies 13
2.3.1 Conventional Detection-Based Architectures 13
2.3.2 Advanced Fusion-Oriented Architectures 14
2.3.3 Segmentation-Oriented Architectures 14
Comparative Summary of Wildfire Datasets and Models 15
2.5 Discussion and Research Implications 15
2.6 Conclusion 16
2.7 Links 16
2.8 Papers Comparison 17
..... 18
Chapter Three 19
DataSets OverView & Comparison 19
3.1 RGB-T 3M 19
3.1.1 Size & Distribution. 19
3.1.2 Modalities & Annotation. 19
3.1.3 Dataset Specifics 19
3.1.4 Use-Cases & Research Impact 20
3.1.5 Remarks & Limitations 20

3.1.6 Overview 20
3.1.7 Reference Link 21
3.2 FireMan-UAV-RGBT 21
3.2.1 Size and Distribution 21
3.2.2 Modality & Annotations 21
3.2.3 Dataset Specifics 22
3.2.4 Use Cases and Research Impact 22
3.2.5 Remarks and Limitations 22
3.2.6 Reference Link 22
3.3 DataSet Comparison 23
Chapter Four 24
Thermal- RGB Fusion for UAV-based Wildfire Monitoring 24
4.1 U-Net 24
4.1.1 Model / Architecture 24
4.1.2 Loss / Optimization 25
4.1.3 Training 25
4.1.4 Metrics 26
4.1.5 Workflow 26
4.2 YOLOv8-ns 33
4.2.1 Model / Architecture 33
4.2.2 Loss / Optimization 34
4.2.3 Training 34
4.2.4 Metrics / Evaluation 35
4.2.5 Data / Input 35
4.2.6 Workflow 36
4.2.7 Model Visualization 38
4.2.8 Learning Curves & Scores 39
4.3 Attention U-Net 40
4.3.1 Model / Architecture 40
4.3.2 Loss / Optimization 41
4.3.3 Training 41
4.3.4 Metrics 42
4.3.5 Workflow 42

The model is trained using a combined loss function consisting of Focal Loss and Dice Loss to handle class imbalance and improve overlap accuracy. Optimization is performed using the Adam optimizer with a learning rate of 5×10^{-4}. 44

Training runs for up to 100 epochs with early stopping applied if validation Dice score does not improve for 15 epochs. The best-performing model is saved based on the highest validation Dice score. 44

- Loaded best saved model 44
- Computed confusion matrix and all metrics 44
- Plotted training & validation loss, Dice, IoU curves 44
- Displayed sample inputs, ground truths, and predicted masks Visual Placeholder: 44
- Learning curves (loss, Dice, IoU) 44
- Prediction comparison (input, ground truth, prediction) 44

4.3.6 Model Visualization 45

4.4 Models Analysis & Discussion 48

1. Comparative Analysis 48
2. Most Important Metrics for Fire Applications 48
3. Final Conclusion 48
4. How to Improve Results: 49

4.4 Data Splitting Strategy 49

4.5 Input Channel Configuration 50

4.6 Input-Level Thermal–RGB Fusion 51

4.6.1 Why Thermal–RGB Fusion Is Used 51

4.6.2 Type of Fusion Employed 51

4.6.3 How Thermal–RGB Fusion Is Applied 52

4.6.4 Comparison with Reference Studies 52

4.6.5 Summery 53

Conclusion 54

References 55

Appendix(1) 56

Appendix(2) 57

Terminology Table 57

Abstract

Wildfires represent a catastrophic threat to ecosystems, necessitating advanced real-time monitoring. This project proposes AI-FireGuard, an innovative framework for wildfire detection using UAVs equipped with dual Thermal and RGB sensors. To address challenges like low visibility and smoke, we evaluated three Deep Learning architectures: U-Net, Attention U-Net, and YOLOv8, utilizing a Thermal-RGB fusion strategy.

Our findings reveal that the U-Net model achieved the highest Precision and IoU, making it the most reliable for precise fire boundary mapping. The Attention U-Net demonstrated superior Recall, effectively identifying obscured fire pixels that standard models might miss. For real-time deployment, YOLOv8 delivered high mAP scores with rapid inference speeds, confirming its suitability for immediate aerial alerts. Overall, the fusion-based approach significantly outperformed single-modality systems, providing a robust and highly accurate solution for early-stage wildfire suppression and disaster management.
Keywords: Wildfire Detection, Thermal-RGB Fusion, U-Net, YOLOv8, UAV Monitoring, Computer Vision.

تمهيد

تمثل حرائق الغابات تهديداً كبيراً للنظم البيئية، مما يستدعي حلولاً متطورة للمراقبة الفورية. يقدم هذا المشروع نظام AI-FireGuard، وهو إطار عمل مبتكر لاكتشاف الحرائق باستخدام الطائرات بدون طيار المزودة بمستشعرات حرارية وبصرية. للتغلب على عوائق الرؤية والدخان، قمنا بتقييم ثلاث بنى للتعلم العميق: U-Net و Attention U-Net و YOLOv8، معتمدين على استراتيجية دمج البيانات (Thermal-RGB Fusion).

كشفت نتائج الدراسة أن نموذج U-Net حقق أعلى مستويات الدقة (Precision) وتقاطع الاتحاد (IoU)، مما يجعله الأكثر موثوقية في رسم خرائط حدود الحريق بدقة. وأظهر نموذج Attention U-Net تفوقاً في الاستدعاء (Recall)، حيث نجح في تحديد بكسلات الحريق المخفية التي قد تتجاهلها النماذج التقليدية. أما بالنسبة للتطبيق الميداني، فقد حقق YOLOv8 نتائج mAP مرتفعة مع سرعة معالجة فورية، مما يؤكد جاهزيته لإرسال التنبيهات الجوية اللحظية. بشكل عام، أثبتت النتائج أن نهج دمج البيانات يتفوق بوضوح على الأنظمة أحادية المصدر، مما يوفر حلاً قوياً وعالي الدقة لإدارة الكوارث وإخماد الحرائق في مراحلها الأولى.

الكلمات المفتاحية: اكتشاف حرائق الغابات، دمج البيانات الحرارية والبصرية، U-Net، YOLOv8، طائرات بدون طيار، الرؤية الحاسوبية.

Chapter One

Introduction

Abstract

Wildfires are among the most destructive natural hazards, posing severe threats to ecosystems, infrastructure, and human life. Early detection and accurate monitoring are essential for reducing damage and enabling timely emergency response. Conventional wildfire monitoring approaches, such as satellite-based or ground sensor systems, often suffer from limited spatial resolution, delayed data availability, and restricted coverage, particularly in remote or complex terrains.

Unmanned Aerial Vehicles (UAVs) equipped with RGB and thermal sensors provide a flexible and high-resolution platform for real-time wildfire surveillance under diverse conditions. In this study, we utilize multi-modal UAV imagery captured during both daytime and nighttime operations, encompassing a wide range of environmental scenarios, including varying illumination, dense smoke, active flames, smoldering regions, and complex backgrounds such as forests and mountainous terrain.

To effectively exploit the complementary nature of the two modalities, thermal–RGB data fusion is employed: thermal imagery emphasizes heat signatures and active fire regions, while RGB imagery provides rich visual and structural context for accurate scene interpretation. We develop and evaluate deep learning models for wildfire detection and segmentation, including U-Net and Attention U-Net for pixel-level fire and smoke segmentation, and YOLOv8 for fast and reliable object-level fire detection. The Attention U-Net further enhances performance by focusing on salient fire-related features while suppressing background noise.

Experimental results demonstrate that the fusion of RGB and thermal imagery significantly improves detection accuracy and robustness compared to single-modality approaches, particularly under low-light, nighttime, smoky, and visually cluttered conditions. The proposed UAV-based multi-modal framework enables accurate, real-time wildfire monitoring and shows strong potential for deployment in early warning systems and emergency response applications.

1.2 Causes

- Wildfires can originate from natural and human-induced factors.
- Lightning strikes are a major natural cause, igniting dry vegetation in remote areas.
- Extended droughts and high temperatures increase fire susceptibility by drying out fuels.
- Human activities, including unattended campfires, discarded cigarettes, and agricultural burning, contribute significantly.
- Power lines and industrial operations can also spark fires in vulnerable regions.
- Wind patterns and terrain slope influence fire spread, complicating detection and control.

1.3 Results

- Our models demonstrate that thermal-RGB fusion significantly improves fire detection accuracy over single-modality data.
- U-Net and Attention U-Net produce detailed segmentation masks, accurately highlighting fire and smoke regions.

- Attention mechanisms enhance focus on small or partially obscured fire regions.
- YOLOv8 enables real-time detection with high precision, suitable for rapid UAV deployment.
- Overall, integrating thermal and RGB channels enables scalable, real-time wildfire monitoring, providing actionable information for emergency response teams.

1.4 Fact

Compared to traditional wildfire monitoring, which relies on satellites or ground sensors with delayed updates and limited resolution, AI-powered systems offer a huge enhancement in early detection. By analyzing thermal and RGB UAV imagery, deep learning models can identify subtle fire and smoke signals in near real-time, even under challenging conditions like dense smoke or complex terrain. This leads to faster alerts, quicker emergency response, and significantly reduced environmental and property damage, demonstrating a clear advantage over conventional approaches.

Chapter Two

Literature Reviews

Recent advances in UAV-based wildfire detection have been driven largely by the availability of RGB–Thermal (RGB-T) datasets and the development of multi-modal deep learning architectures. Existing works can be clustered according to (1) dataset characteristics, (2) fusion strategy and model architecture, and (3) detection or segmentation performance.

2.1 Introduction

The effectiveness of UAV-based wildfire detection systems is strongly influenced by the quality, diversity, and modality of the datasets used for training and evaluation. Recent research has increasingly focused on multi-modal RGB–Thermal (RGB-T) datasets, as they provide complementary visual and thermal information that improves detection robustness under challenging environmental conditions such as smoke, low illumination, and complex vegetation. This section provides a detailed comparative analysis of prominent wildfire datasets and associated methodologies, clustering them according to dataset type, model architecture, fusion strategy, and experimental performance.

2.2 Dataset Classification Based on Modalities and Acquisition Conditions

2.2.1 RGB–Thermal UAV Datasets

2.2.1.1 RGBT-3M Dataset

The RGBT-3M dataset represents one of the most comprehensive RGB–Thermal UAV datasets for wildfire detection. It consists of paired RGB and

thermal images captured during both daytime and nighttime operations, covering a wide range of real-world wildfire scenarios. The dataset includes scenes with dense smoke, varying fire intensities, poor lighting conditions, and complex terrain backgrounds, such as forests and mountainous regions.

To improve robustness, the dataset incorporates synthetic smoke simulation, enabling the models to generalize better to severe visibility degradation. This dataset is primarily designed for object-level fire detection, with annotations compatible with modern one-stage detectors.

2.2.1.2 FireMan-UAV-RGBT Dataset

The FireMan-UAV-RGBT dataset is a RGB–Thermal video dataset captured during controlled wildfire experiments in Finnish boreal forests using DJI UAV platforms. The data were collected at altitudes below 120 meters, ensuring high spatial resolution and realistic fire dynamics.

A distinctive feature of this dataset is its semi-automatic annotation process, where thermal imagery is used to guide the labeling of RGB frames. This approach reduces annotation ambiguity and improves consistency in fire region labeling. Compared to RGBT-3M, FireMan-UAV-RGBT focuses on a specific ecological environment, making it particularly relevant for boreal forest fire monitoring.

2.2.2 RGB-Only UAV Datasets

2.2.2.1 U3UNet Dataset

The dataset used in the U3UNet study consists of high-resolution RGB UAV images focusing on pixel-level segmentation of fire regions. Unlike RGB–Thermal datasets, this dataset relies solely on visible-spectrum imagery and is primarily designed to capture complex flame shapes and fire boundaries.

Although effective under good illumination conditions, RGB-only datasets remain sensitive to smoke occlusion and low-light environments, limiting their applicability in nighttime or early-stage fire detection.

2.3 Model Architectures and Fusion Strategies

2.3.1 Conventional Detection-Based Architectures

2.3.1.1 YOLOv8-Based Detection on RGBT-3M

YOLOv8 was employed on the RGBT-3M dataset due to its strong balance between detection accuracy and inference speed. Feature-level fusion was applied to integrate RGB and thermal information, allowing the network to exploit thermal heat signatures while retaining spatial and contextual details from RGB images.

The model achieved a mean Average Precision (mAP@0.5) of 96.3%, with a processing speed of approximately 39 frames per second, demonstrating suitability for near real-time wildfire monitoring.

2.3.1.2 YOLOv8 and ResNet50 on FireMan-UAV-RGBT

In the FireMan-UAV-RGBT study, both YOLOv8 and ResNet50 architectures were evaluated using unimodal (RGB-only, thermal-only) and multimodal (RGB-T) inputs. Experimental results consistently showed that RGB–Thermal fusion outperformed unimodal configurations, confirming the effectiveness of multi-modal sensing for wildfire detection.

However, cross-dataset evaluation revealed reduced generalization performance, indicating that models trained on region-specific datasets may struggle when deployed in different ecological or climatic environments.

2.3.2 Advanced Fusion-Oriented Architectures

2.3.2.1 MCDet: Target-Aware Fusion Framework

The MCDet framework introduces a Target-Aware Fusion strategy to address the ambiguity and instability inherent in RGB–Thermal data. The proposed Multidimensional Representation Collaborative Fusion (MRCF) module models global relationships between RGB and thermal features using a state-space formulation, while Deformable Convolutions enhance sensitivity to small or partially occluded fire regions.

Additionally, the Content-Guided Attention Network (CGAN) dynamically weights modality contributions based on scene context, suppressing misleading RGB reflections or non-fire thermal hotspots. The use of Weighted IoU (WIoU) loss further improves discrimination between true fire regions and false positives.

Despite achieving a moderate mean detection accuracy of 77.5%, MCDet demonstrated superior robustness under dense smoke, vegetation clutter, and low-light conditions.

2.3.3 Segmentation-Oriented Architectures

2.3.3.1 U3UNet for Fire Region Segmentation

U3UNet is a specialized U-Net-based architecture designed for precise segmentation of wildfire regions in UAV imagery. Its design focuses on preserving fine-grained spatial details and accurately capturing irregular fire boundaries.

While U3UNet delivers high segmentation accuracy under favorable conditions, its reliance on RGB imagery limits robustness in scenarios with

heavy smoke or nighttime conditions, where thermal information would provide critical advantages.

Comparative Summary of Wildfire Datasets and Models

Dataset / Model	Task	Modality	Key Metrics	Notable Outcomes
RGBT-3M + YOLOv8	Detection	RGB-T	mAP@0.5 = 96.3%	Real-time performance; strong robustness under dense smoke
MCDet	Detection	RGB-T	Accuracy = 77.5%	Target-aware fusion improves stability in challenging conditions
FireMan-UAV-RGBT	Detection	RGB-T	High dataset-specific accuracy	Limited cross-dataset generalization
U3UNet	Segmentation	RGB	High pixel-level accuracy	Reduced reliability under nighttime conditions

2.5 Discussion and Research Implications

The comparative analysis highlights several key observations:

Multi-modal RGB–Thermal datasets significantly enhance detection reliability, particularly in low-light and smoke-obscured environments.

Large-scale datasets favor high detection accuracy and real-time deployment, while fusion-centric architectures prioritize robustness under complex conditions.

Geographic and environmental specificity remains a major limitation, emphasizing the need for globally diverse datasets.

Temporal modeling and fire progression analysis are largely absent and represent an important direction for future research.

2.6 Conclusion

In conclusion, existing UAV-based wildfire datasets and detection models demonstrate the clear advantages of RGB–Thermal fusion for robust wildfire monitoring. Datasets such as RGBT-3M and FireMan-UAV-RGBT provide valuable benchmarks, while architectures like MCDet and U3UNet highlight different trade-offs between accuracy, robustness, and segmentation precision. These findings strongly motivate the development of diverse, multi-modal, and temporally-aware datasets, which will be essential for next-generation wildfire detection systems.

2.7 Links

https://www.mdpi.com/2072-4292/17/15/2593

https://www.mdpi.com/2072-4292/17/15/2593

https://ieeexplore.ieee.org/document/10710657

https://ieeexplore.ieee.org/document/10944154

2.8 Papers Comparison

Aspect	RGBT-3M	MCDet	UAV-RGBT	Real-time UAV smoke segmentation.
Full Title	RGBT-3M: A Large-Scale Multi-Modal Benchmark for RGB-T Image Fusion and Fire Detection	MCDet: Target-Aware Fusion for RGB-T Fire Detection	Fireman -UAV-RGBT: a novel UAV-based RGB-Thermal video dataset for the detection of wildfires in the Finnish forests.	Detecting Wildfires on UAVs with Real-Time Segmentation Trained by Larger Teacher Models
Publication Year	2024 (used in 2025 studies)	2025	AUG 2023.	2024
Publisher	MDPI – Remote Sensing Journal	Elsevier – Fire Safety Journal	Zenodo.	arXiv.
Scientific Goal	To build a large-scale benchmark dataset for evaluating RGB-T image fusion models in fire and outdoor detection tasks.	To develop a deep learning model for effective RGB-T image fusion improving real-time wildfire detection accuracy.	To create and evaluate a new UAV-based RGB-Thermal video dataset for improving wildfire detection in boreal forests using deep learning models.	Achieve accurate onboard smoke segmentation without requiring pixel-level human annotations.
Study Focus	Dataset creation and benchmarking framework.	Algorithmic innovation and model performance enhancement.	Advancing wildfire detection through a novel UAV-based RGB-Thermal dataset and deep learning evaluation.	Bounding-box–based weak supervision, pseudo-label generation using SAM and BoxSnake, and teacher–student distillation to train lightweight PIDNet models for UAV deployment.
Data Type	Paired RGB and Thermal images from ground and UAV perspectives.	RGB-T video frames captured from UAVs in dynamic outdoor wildfire environments.	RGB videos Thermal videos Paired RGB-Thermal videos.	RGB UAV imagery with bounding-box annotations; pseudo-label masks generated automatically.
Dataset Size	Over 285,000 image pairs covering day/night and various illumination conditions.	Approx. 32,000 video frames and 12,000 selected wildfire scenes.	33 RGB videos+10 RGB-Thermal paired videos of total 104,816 frame.	Several thousand images from AI for Mankind v1.0/v2.0 and the Finnish UAV smoke dataset (controlled burns + smokeless negative samples).
Fusion Technique	Feature-level and decision-level fusion using CNN and Transformer models.	Novel MRCF (Multidimensional Representation Collaborative Fusion) method.	Early fusion by concatenating RGB and Thermal images into a single 6-channel input.	Pseudo-label fusion from SAM and BoxSnake-generated masks combined to guide student model training.
Base Model	ResNet, ViT, Swin Transformer, DenseNet (for benchmarking).	Custom YOLOv8n-based model with Target-Aware Fusion Module (MCDet).	ResNet50 and YOLOv8.	PIDNet (S, M, L variants) used as the lightweight real-time segmentation model.
Training Approach	Unified training and evaluation across all models for fair benchmarking.	Attention-based selective training focused on small and early fire targets.	Transfer learning by initializing models with pre-trained ImageNet weights and fine-tuning them.	Weak supervision using bounding boxes, pseudo-label generation, and teacher–student knowledge distillation.

Results Summary	RGB-T fusion improved F1-Score by +9.8% over RGB-only baselines.	MCDet achieved +13.6% better accuracy for early-stage wildfire detection.	Internal test accuracy reached 100% but generalization to external datasets remained challenging.	PIDNet achieves strong segmentation, runs in real time (~25 fps) on Jetson Orin NX, and detects smoke up to 9.7 km. mIoU 63.3%
Scientific Contribution	First large-scale benchmark for RGB-T fusion in fire and rescue contexts.	Introduced a novel target-aware fusion algorithm for real-time UAV deployment.	A novel RGB-Thermal wildfire dataset and benchmark for boreal forests.	Demonstrates viable real-time smoke segmentation on UAVs using weak labels instead of costly pixel-level masks.
Technical Innovation	1) Large-scale dataset creation. 2) Analysis of illumination and terrain effects.	1) Target-aware fusion module. 2) Robustness under smoke and variable lighting.	Early fusion of UAV-captured RGB-Thermal video for wildfire detection.	Integrates pseudo-labeling (SAM + BoxSnake) with distillation to train lightweight models for onboard UAV deployment.
Applications	Benchmarking and evaluation of multimodal AI systems.	Operational UAV wildfire monitoring and prevention systems.	Wildfire detection and monitoring in boreal forests using drones.	UAV wildfire monitoring and real-time smoke detection
Best Use Case	Data analysis, model benchmarking, simulation projects.	Applied AI development and real-time fire detection systems.	Early detection of wildfires in remote boreal forests using drone-based monitoring.	Deploying lightweight models on UAVs for early smoke detection and alerts.
Limitations	Lacks temporal modeling (no video or sequence-based context).	Limited by dataset size and environmental diversity.	1)Limited generalization 2)Controlled environment 3)Manual registration.	Limited dataset diversity, pseudo-label noise, and false positives in smokeless scenes.
Academic Recommendation	Best suited for research and dataset benchmarking studies.	Best suited for real-time and practical AI field systems.	Best suited for early case detection of wildfires.	Recommended for researchers and practitioners in UAV-based wildfire detection and edge AI deployment.

Chapter Three

DataSets OverView & Comparison

3.1 RGB-T 3M

3.1.1 Size & Distribution

This dataset, hosted by the University of Science and Technology of China (USTC), includes approximately **1,367** RGB–Thermal paired images. It is referenced in wildfire detection research as a compact but valuable benchmark dataset for multimodal fire analysis.

3.1.2 Modalities & Annotation

The dataset provides paired RGB and thermal images, captured under varied illumination and fire intensity conditions. Annotations support semantic segmentation tasks, labeling fire and non-fire regions, enabling multimodal training for deep learning models.

3.1.3 Dataset Specifics

- Format: Image pairs (RGB + Thermal)
- Approximately total samples: 1,367
- Target task: Fire detection and segmentation
- Annotation type: Pixel-level segmentation masks
- Data balance: Fire vs non-fire, collected in both day and night settings

3.1.4 Use-Cases & Research Impact

It supports tasks such as wildfire detection, multimodal fusion learning, and real-time UAV-based thermal imaging analysis. Commonly used in comparative studies alongside datasets like FLAME-1 and FireMan-UAV-RGBT.

3.1.5 Remarks & Limitations

The dataset's modest scale limits generalization but offers a valuable baseline for fusion-based fire detection. Access may require university credentials or approval.

3.1.6 Overview

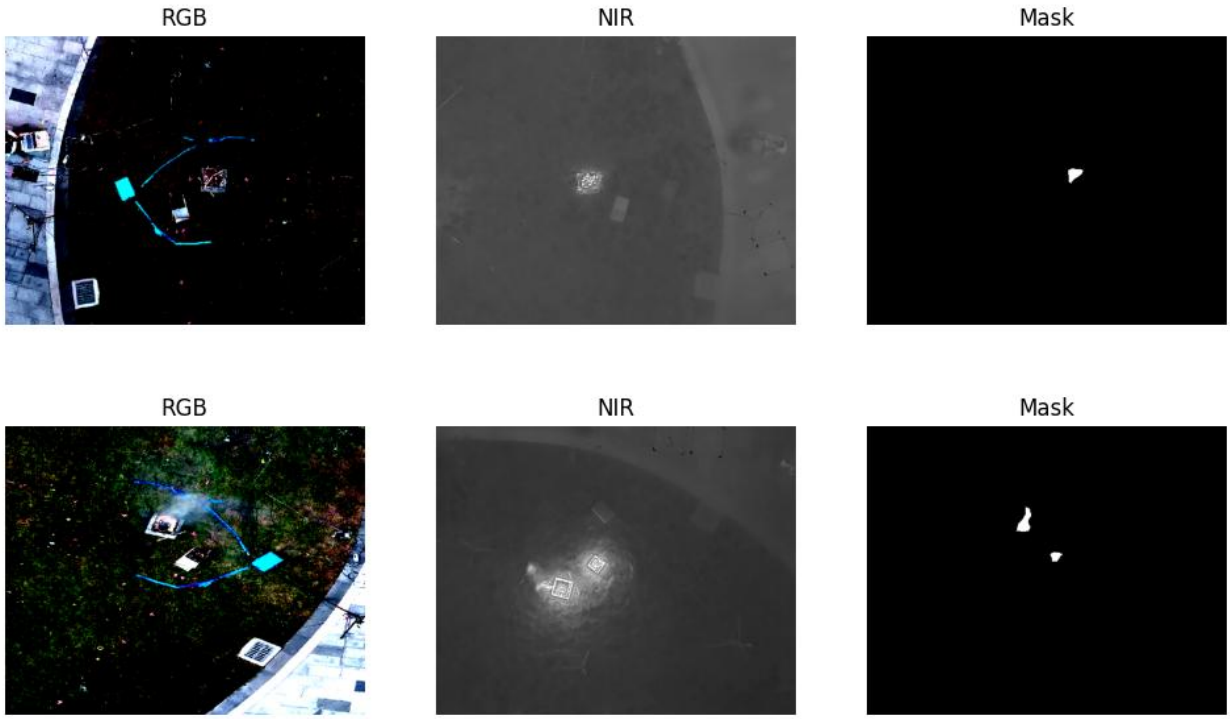

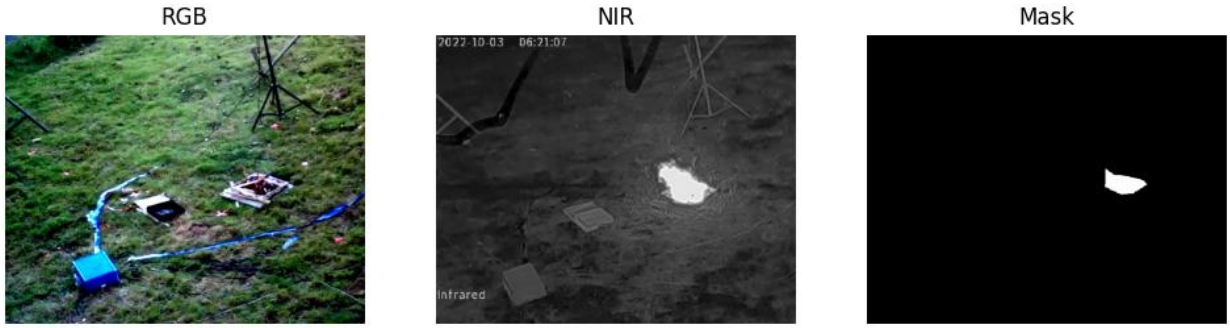

3.1.7 Reference Link

https://rec.ustc.edu.cn/share/75aae550-8d26-11ee-98a1-ed4fd857523e

3.2 FireMan-UAV-RGBT

3.2.1 Size and Distribution

Total dataset size: ~39.4 GB.

34 RGB videos + 20 paired RGB-Thermal videos.

Captured across four controlled burn sites in Finnish boreal forests.

Covers varied flight paths, altitudes, and scene complexity.

3.2.2 Modality & Annotations

Modality: RGB + Thermal synchronized video.

Annotations: YOLO-format bounding boxes for fire and related objects.

Both manual and semi-automatic annotation methods applied.

3.2.3 Dataset Specifics

Video resolution: High-resolution UAV footage.

Multi-environment coverage: Forests, vegetation, varied lighting conditions.

Includes sample frames and labeling reference files for model training.

3.2.4 Use Cases and Research Impact

Primary use: Training and benchmarking UAV-based wildfire detection models.

Supports multimodal fusion research (RGB + Thermal).

Enables development of early wildfire detection systems, environmental monitoring, and UAV operational testing.

3.2.5 Remarks and Limitations

Strengths: Real UAV footage, synchronized RGB-Thermal, YOLO-compatible annotations.

Limitations: Data comes from controlled burns, so real wildfire variability may be underrepresented.

Dataset is moderate in size compared to large benchmarks but rich in temporal video sequences.

3.2.6 Reference Link

Zenodo dataset: https://zenodo.org/records/13732947

IEEE Paper: https://ieeexplore.ieee.org/document/10710657

3.3 DataSet Comparison

Dataset Name	Paper Date	Tools & Techniques Used	Size & Data Description
RGBT-3M	2025	RGB-T fusion networks, deep learning, UAV-based fire detection, multi-scenario evaluation	3M paired RGB-T images, multi-scenario (forest, urban, grassland) UAV dataset
MCDet	2025	Target-aware RGB-T fusion, transformer-based feature alignment, UAV fire detection	25K RGB-T image pairs collected in diverse lighting and thermal conditions
FireMan-UAV-RGBT	2024	Deep CNNs, spatiotemporal fusion, UAV video-based wildfire detection	High-resolution RGB-T video dataset covering various wildfire scenarios
UAVs-FFDB	2023	YOLOv5, dataset augmentation, UAV image-based detection pipeline	10K RGB and IR images from UAV flights over forest fire zones
FLAME	2022	Machine learning models, multispectral analysis, fire luminosity estimation	Airborne images of fires with brightness and heat annotations
UAV-StrawFire	2021	Real-time CNN-based fire monitoring, UAV RGB+IR fusion techniques	Dataset focusing on straw fire monitoring under controlled field experiments

#Images / #Videos	Annotations	Input	Output	Citations
≈3,000,000 images	Bounding boxes, fire masks, scene categories	RGB and Thermal images (UAV)	Fire detection and localization	New (few citations)
≈25,000 images	Bounding boxes, fire intensity levels	RGB and Thermal images (UAV)	Enhanced fire detection via fusion model	New (few citations)
≈200 videos (RGB-T pairs)	Frame-level fire presence, temporal labels	RGB-T video frames (UAV)	Wildfire detection and video-level classification	Moderate (10–20 citations)
≈10,000 images	Bounding boxes, fire/no-fire labels	RGB and IR images (UAV)	Forest fire presence detection	20+ citations
≈5,000 images	Fire luminosity and temperature metrics	Airborne RGB and thermal imagery	Fire intensity and luminosity estimation	40+ citations
≈2,000 paired images and short video clips	Bounding boxes, fire pixel-level segmentation	Visible and infrared UAV images	Real-time straw fire monitoring	30+ citations

Chapter Four

Thermal- RGB Fusion for UAV-based Wildfire Monitoring

This chapter provides an explanation of the training process steps and the results achieved, in addition to the architecture of the models and their evaluation in terms of effectiveness and accuracy.

4.1 U-Net

U-Net is widely used in fire detection and mask segmentation because it enables precise pixel-level localization of fire and smoke regions. Its encoder–decoder architecture with skip connections preserves fine spatial details while capturing high-level contextual features. This design is particularly effective for irregular fire shapes and fuzzy smoke boundaries. U-Net performs well even with limited labeled data, which is common in wildfire datasets. Consequently, it is well suited for accurate segmentation in complex UAV imagery.

4.1.1 Model / Architecture

Model type: U-Net (custom implementation with 4 downsampling and upsampling layers)

Input channels: in_ch = 4 (RGB + Thermal or other 4-channel input)

Output channels: out_ch = 1 (binary segmentation mask)

DoubleConv filters: [64, 128, 256, 512]

Upsample mode: bilinear

Skip connections: concatenation with corresponding encoder features

4.1.2 Loss / Optimization

The ComboLoss, combining Focal Loss and Dice Loss, is used to balance class imbalance and segmentation accuracy in fire and smoke masks. Focal Loss (α = 0.25, γ = 2) emphasizes hard-to-classify fire pixels while reducing the impact of easy background samples. Dice Loss with smoothing (smooth = 1) improves overlap between predicted and ground-truth masks, especially for small or fragmented fire regions. The Adam optimizer is selected for its adaptive learning capability and stable convergence. A learning rate of 5×10^{-4} ensures efficient training while avoiding unstable updates.

- Loss function: ComboLoss = 0.5 * FocalLoss + 0.5 * DiceLoss
- FocalLoss hyperparameters: alpha = 0.25, gamma = 2
- DiceLoss smoothing: smooth = 1
- Optimizer: Adam
- Learning rate: 5e-4

4.1.3 Training

Number of epochs: num_epochs = 100

Early stopping patience: patience = 15 (stop if no improvement in Dice for 15 epochs)

Device: GPU if available (cuda), else CPU

Batch size: Not explicitly defined; depends on train_loader

Binary threshold: 0.5 (sigmoid output)

4.1.4 Metrics

Dice, IoU, Accuracy, Precision, Recall, F1-score:

The selected evaluation metrics provide a comprehensive assessment of segmentation performance. Dice coefficient and Intersection over Union (IoU) measure the spatial overlap between predicted and ground-truth fire masks. Accuracy evaluates overall pixel classification correctness. Precision and Recall quantify false positives and false negatives, which are critical for avoiding missed fire regions. The F1-score balances precision and recall, offering a robust indicator of segmentation reliability.

4.1.5 Workflow

1. Data Preparation

- Dataset Setup
- Downloaded RGB-T wildfire dataset (ZIP files).
- Mounted Google Drive and extracted the data.
- Dataset structure: rgb/, nir/, mask/.
- Library Installation

-Deep learning: PyTorch, TensorFlow, Keras

- Image processing: OpenCV, scikit-image, imageio

- Data & visualization: numpy, pandas, matplotlib, seaborn, plotly, tqdm

- Augmentation: Albumentations
- Data Loading & Verification

- Loaded RGB, NIR, and mask images.
- Skipped missing/incomplete samples.
- Total valid samples: 1367
- Image shapes: RGB (420,420,3), (512,640,3); NIR & masks (512,640)
- Mask unique values: 0 (background), 255 (wildfire)

2. Data Preprocessing & Augmentation

- Augmentation Strategy (Albumentations)
- Training: horizontal/vertical flips, rotations, color jitter, Gaussian blur, resize, normalization
- Validation: resize and normalization only
- Custom Dataset Class
- Combined RGB + NIR → 4-channel input
- Applied augmentations dynamically
- Converted to PyTorch tensors [B,4,H,W]for images and [B,1,H,W] for masks
- Data Split & Loading
- Train: 80%, Validation: 20%
- Batch size: 4, shuffle enabled for trainingVisual Placeholder:
- Sample RGB/NIR/mask images before and after augmentation

3. Model Design

- Architecture: Custom 4-channel UNet
- Encoder: 4 downsampling blocks
- Bottleneck + decoder with skip connections

- Output: single-channel mask
- Loss Function: ComboLoss = 0.5 * Focal Loss + 0.5 * Dice Loss
- Metrics: Dice coefficient, IoU, Accuracy, Precision, Recall, F1-score

4. Training & Optimization

- - Optimizer: Adam (lr=5e-4)
- - Epochs: 100, Early stopping patience: 15
- - GPU acceleration (CUDA)
- - Training loop:

  1. Forward pass → predictions
  2. Compute loss
  3. Backpropagation & optimizer step
  4. Compute metrics for train & validation

- Early stopping monitored on best Dice coefficient

  Best model saved as best_model.pth

5. Evaluation & Visualization

- Loaded best saved model
- Computed confusion matrix and all metrics
- Plotted training & validation loss, Dice, IoU curves
- Displayed sample inputs, ground truths, and predicted masks Visual Placeholder:
- Learning curves (loss, Dice, IoU)
- Prediction comparison (input, ground truth, prediction)

4.1.5 Model Visualization

4.1.5.1 Architecture Visualize

- Input (4 channels)
- Down path: DoubleConv blocks with [64, 128, 256, 512]
- MaxPooling for downsampling
- Up path: Upsample + DoubleConv blocks [512→256→128→64]
- Skip connections from encoder to decoder
- Output: 1 channel (binary mask)

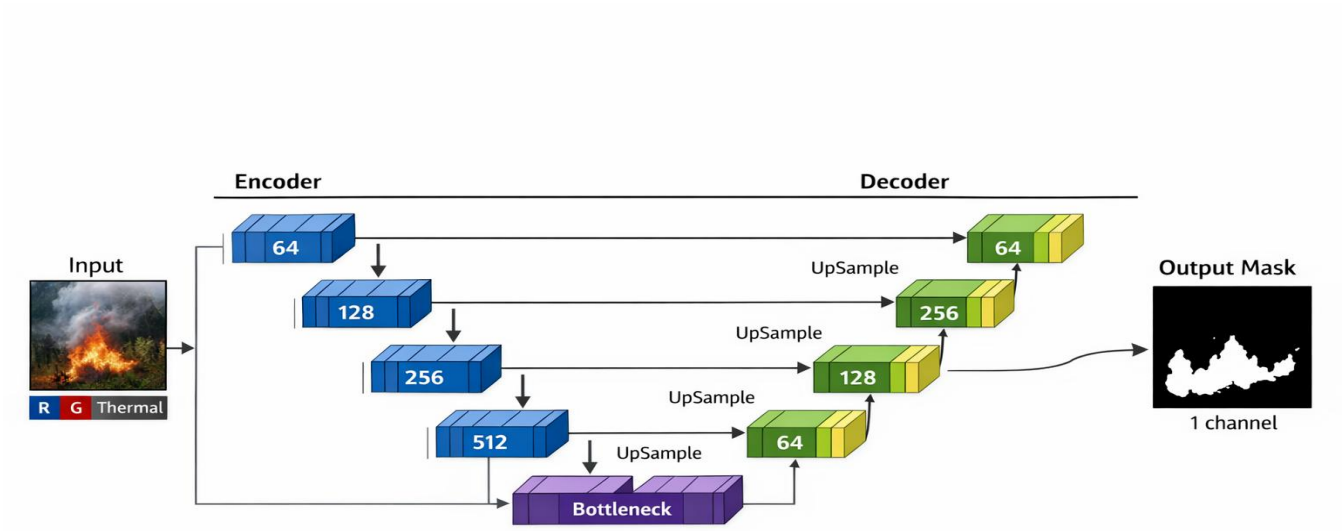

4.1.5.2 Ground Truth vs Predicted Mask

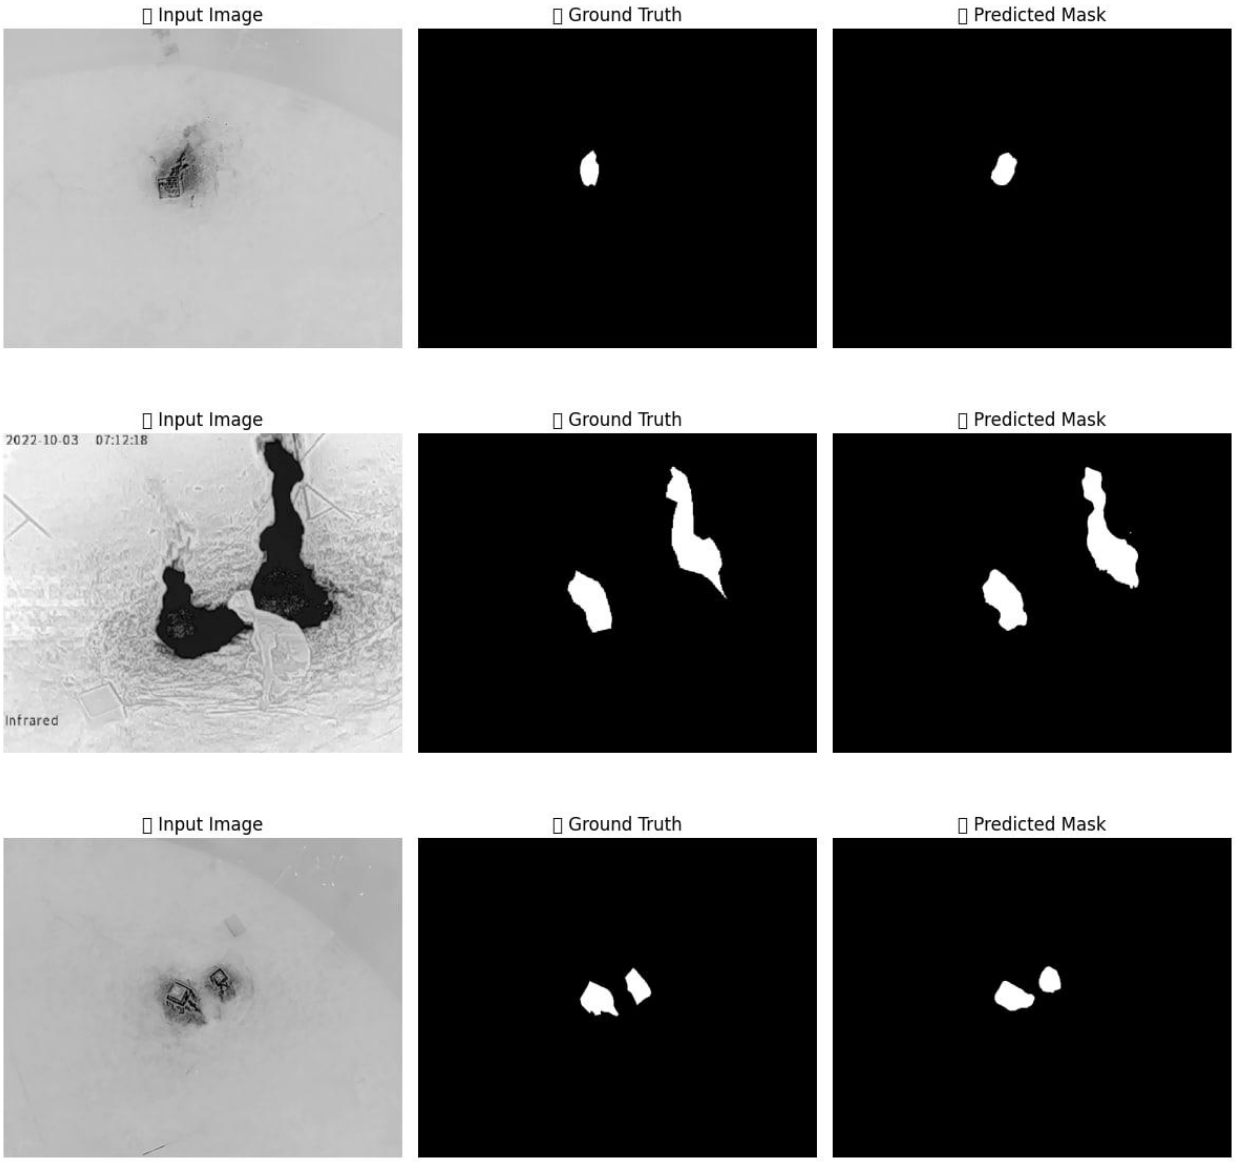

4.1.5.3 Learning Curves ,Confusion Matrix & Scores

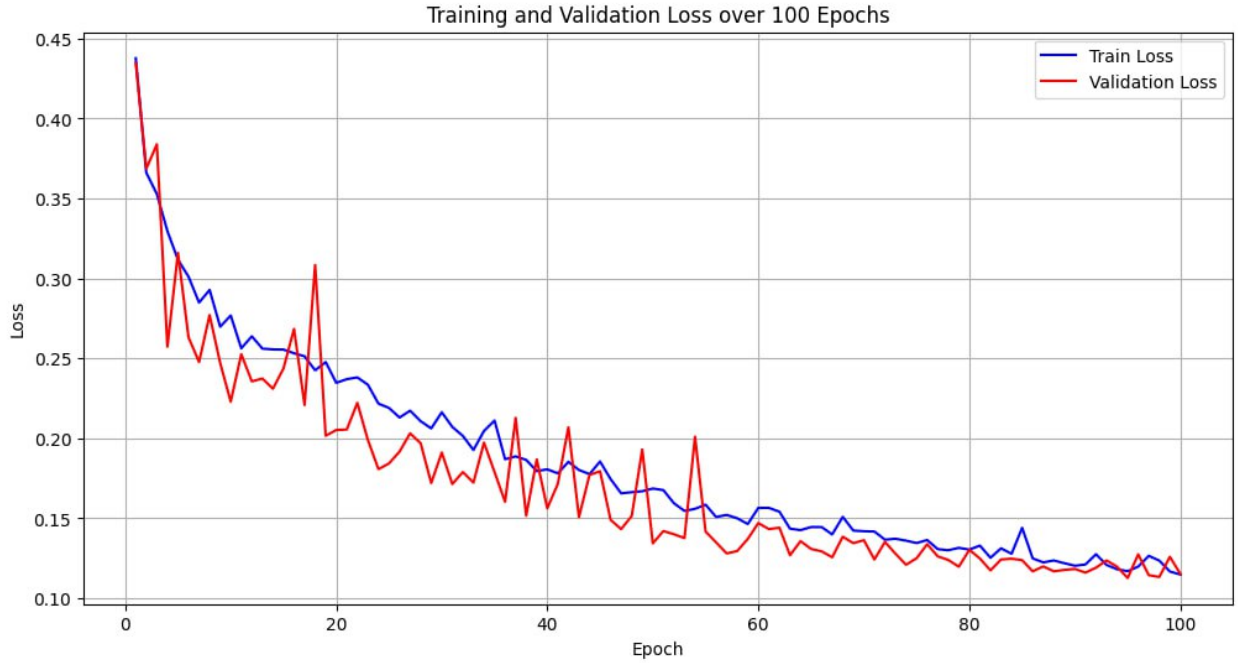

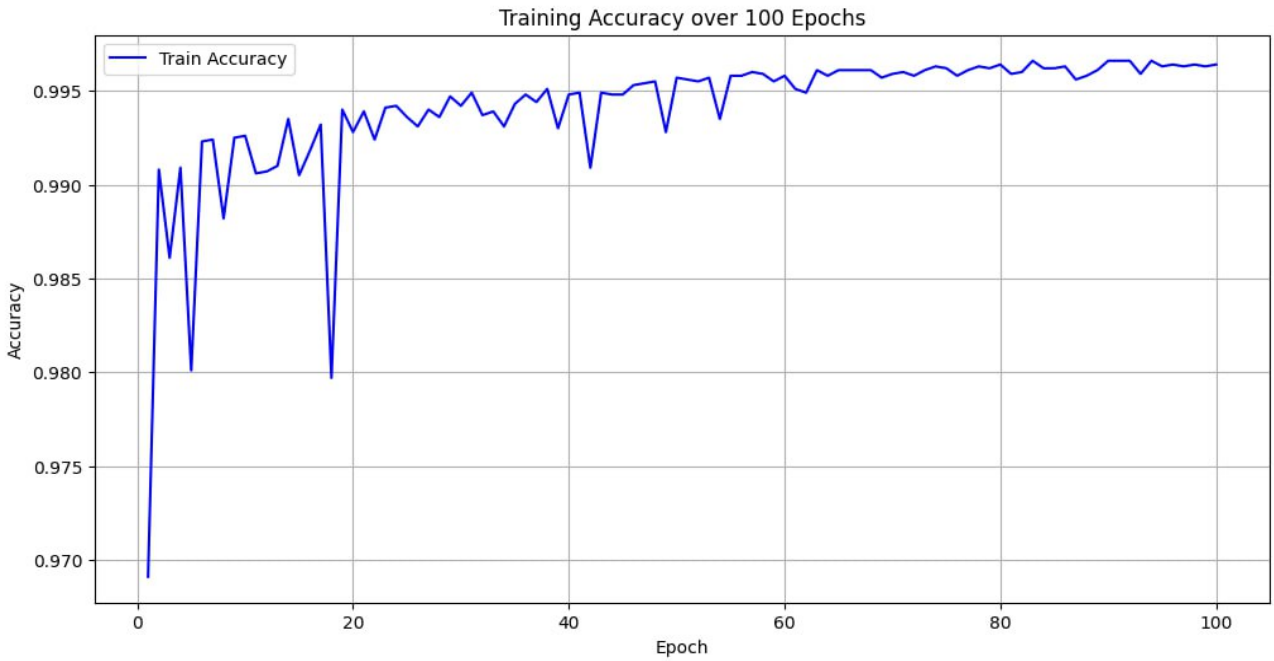

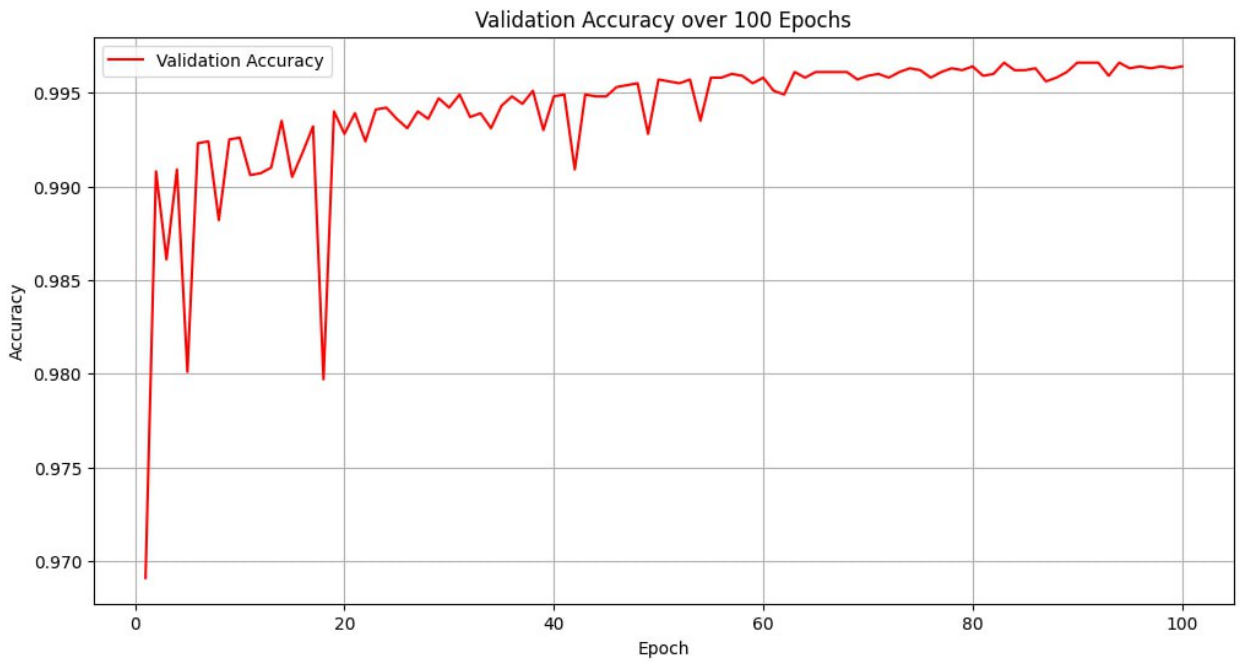

Model	Accuracy	Precision	Recall	F1-score	IoU
U-Net	0.9965	0.7583	0.8494	0.8013	0.6684

4.2 YOLOv8-ns

YOLOv8-Nano is used in fire detection and mask segmentation due to its lightweight architecture and fast inference speed, making it suitable for real-time UAV deployment. Despite its small size, it maintains competitive detection accuracy for fire and smoke regions. Its low computational cost allows deployment on edge devices with limited resources. YOLOv8-Nano supports joint detection and segmentation tasks, enabling efficient fire localization and mask generation. This makes it ideal for early wildfire detection scenarios where speed is critical.

4.2.1 Model / Architecture

Model type: YOLOv8n (Nano)

Input channels: 3 (RGB images)

Output channels: 1 (single-class detection: fire)

Backbone: CSP / Conv / Bottleneck (YOLOv8 default)

Neck: PAN-FPN for feature aggregation

Head: YOLO detection head (predicts bounding boxes, objectness, class confidence, optional mask)

Mask output: Enabled (overlap_mask=True, mask_ratio=4)

Pretrained: Yes (pretrained=True)

4.2.2 Loss / Optimization

The YOLOv8 training setup employs specialized loss components to optimize detection accuracy and localization quality. CIoU/GIoU box loss (weighted by 7.5) improves bounding-box regression by considering overlap, distance, and aspect ratio. Binary Cross-Entropy classification loss (cls = 0.5) ensures stable fire/non-fire class prediction, while Distribution Focal Loss (dfl = 1.5) refines bounding-box precision. Optimization is handled by YOLOv8's automatic optimizer, using a carefully scheduled learning rate (lr_0 = 0.001 to lrf = 0.01). Momentum, weight decay, and warmup epochs enhance convergence stability and generalization.\

Loss functions:

- Box loss: CIoU / GIoU (YOLOv8 default, weighted by box=7.5)
- Classification loss: BCE with weight cls=0.5
- Distribution Focal Loss (DFL): dfl=1.5
- Optimizer: Automatic (YOLOv8 internal)
- Initial learning rate: lr0=0.001
- Final learning rate: lrf=0.01
- Momentum: 0.937
- Weight decay: 0.0005
- Warmup epochs: 3.0

4.2.3 Training

Number of epochs: 100

Early stopping patience: 25 epochs (patience=25)

Batch size: 16 (batch=16)

Image size: 640 × 640 (imgsz=640)

Deterministic training: True (deterministic=True)

Random seed: 42 (seed=42)

Single class: True (single_cls=True)

Rectangular training: False (rect=False)

Cosine learning rate: False (cos_lr=False)

Dropout: 0.0

4.2.4 Metrics / Evaluation

Metrics calculated: Precision, Recall, mAP@0.5, mAP@0.5–0.95, Box/Cls/DFL losses

For YOLOv8-based fire detection, performance is evaluated using Precision and Recall to measure the correctness and completeness of detected fire regions. mAP@0.5 assesses average precision at a 50% IoU threshold, while mAP@0.5–0.95 provides a stricter evaluation across multiple IoU thresholds, reflecting localization accuracy. Additionally, Box, Classification (Cls), and Distribution Focal Loss (DFL) values are monitored during training to ensure accurate bounding-box regression, class prediction, and precise object localization. These metrics together provide a comprehensive assessment of detection and localization performance.

4.2.5 Data / Input

Training images: 759

Validation images: 189

Fire regions annotated: 2208

Mask support: Enabled via overlap_mask

4.2.6 Workflow

1. Data Preparation & Preprocessing

- Objective: Convert raw dataset into YOLO-compatible format for training.

Steps:

- **Dataset Organization:**
- Images: dataset/images/
- Masks: dataset/masks/
- Labels output: dataset/labels/

Mask Conversion:

- Each mask is converted to YOLO .txt format.
- Multi-class handling using class_map
- Bounding boxes extracted from contours of non-zero regions.
- Minimum area filtering to remove noise.
- Morphological operations applied to refine masks.

Dataset Split:

- Train/Validation split ratio: 80/20.
- YAML file (dataset.yaml) automatically generated for YOLO.

Output:

- .txt label files aligned with each image.
- Cleaned, preprocessed dataset ready for YOLOv8 training.

2. Model Definition

- Objective: Initialize YOLOv8 model for fire detection.
- Details:
- Pretrained model used: yolov8n.pt (nano version for speed).
- Single-class detection (single_cls=True) targeting fire regions.
- Input image size: 640x640 pixels.

3. Training Loop

- **Objective:** Train YOLOv8 model with early stopping and monitoring.
- Key Parameters:
- Epochs: 100
- Batch Size: 16
- Learning Rate: Initial = 0.001, Final = 0.01
- Optimizer: Auto
- Patience: 25 (early stopping)
- Loss weights: Box = 7.5, Classification = 0.5, DFL = 1.5
- Warmup Epochs: 3

4 Loss Functions & Metrics

YOLOv8 Loss Components:

- 1. Box Loss: Measures localization error of predicted bounding boxes.

- 2. Classification Loss: Cross-entropy between predicted class and ground truth.
- 3. DFL Loss: Distribution Focal Loss for bounding box refinement.

Evaluation Metrics:

- Precision: Fraction of correct positive predictions.
- Recall: Fraction of true positives detected.
- mAP50: Mean Average Precision at 0.5 IoU threshold.
- mAP50-95: mAP averaged across IoU thresholds 0.5–0.95.
- F1-score: Harmonic mean of Precision and Recall.

4.2.7 Model Visualization

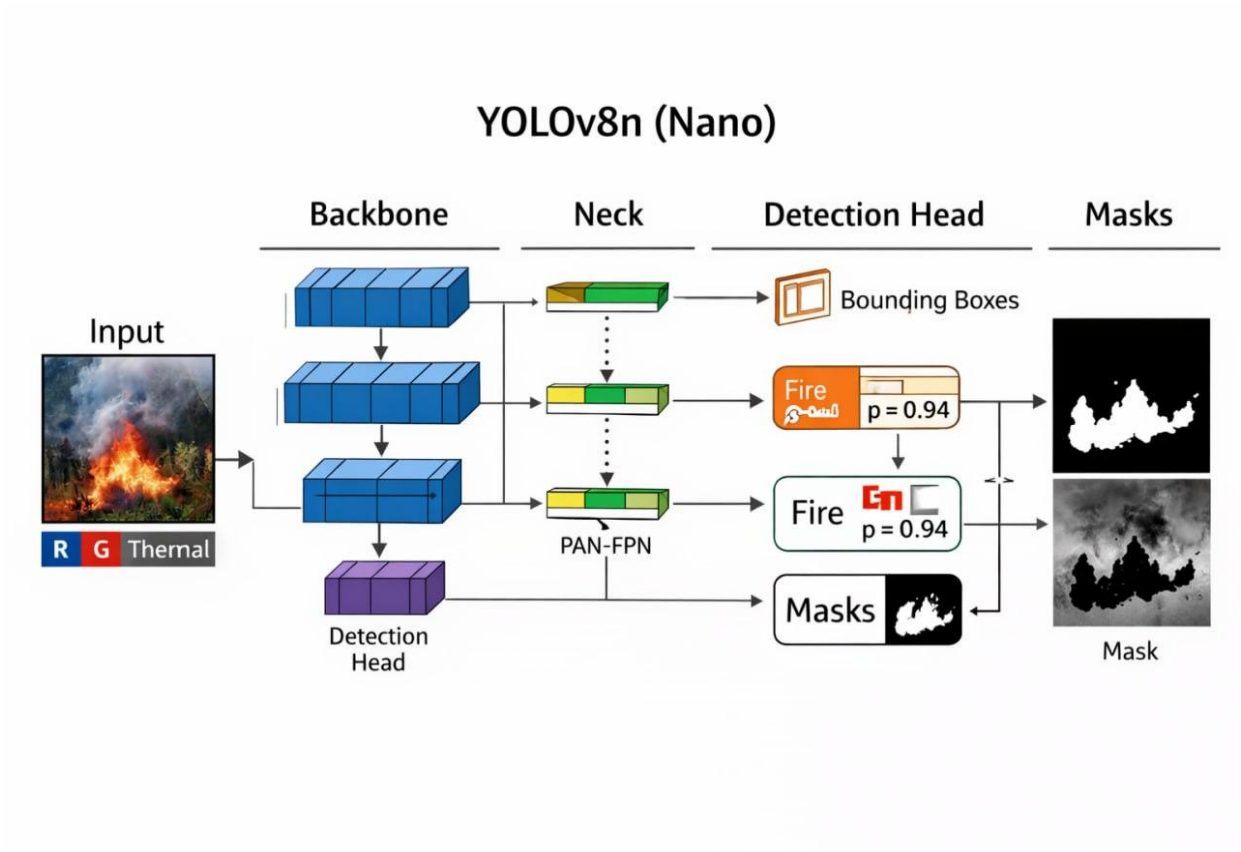

4.2.8 Learning Curves & Scores

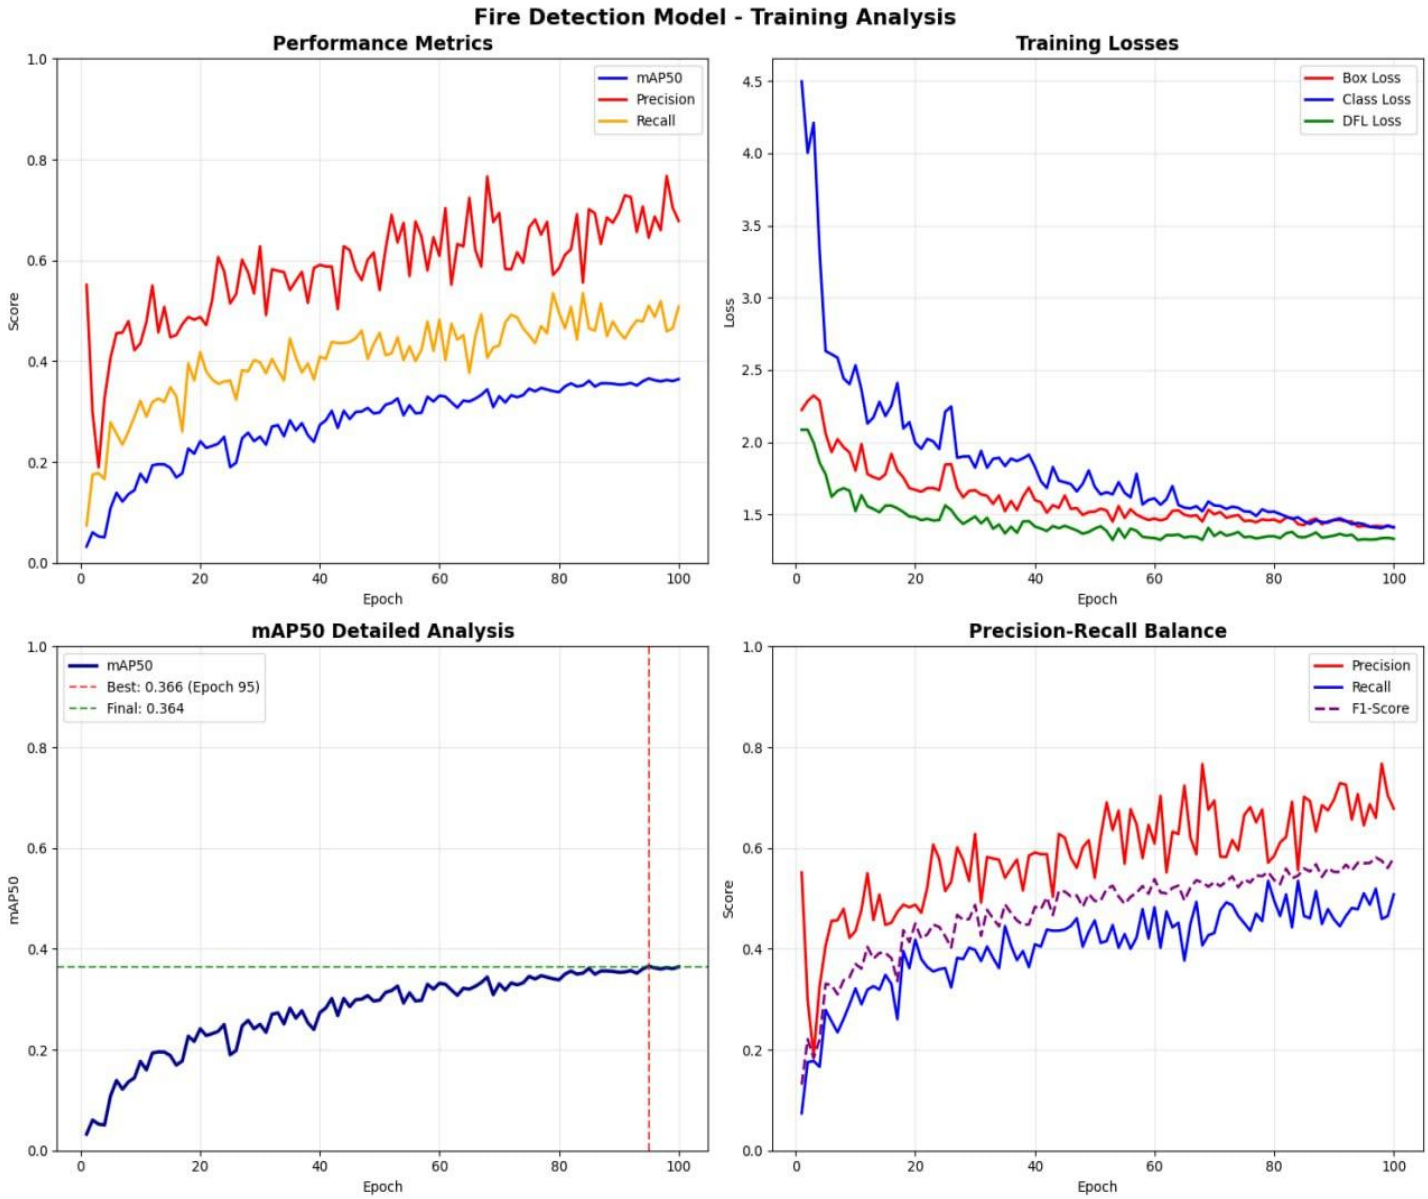

Metric	Initial	Final	Best
Precision	0.5513	0.6781	0.7673
Recall	0.0742	0.5079	0.5348
mAP@50	0.0326	0.3644	0.3657

4.3 Attention U-Net

Attention U-Net is used in fire detection and mask segmentation to focus the model on relevant fire and smoke regions while suppressing background noise. Its attention gates enhance feature selection, which is crucial in complex scenes with vegetation, clouds, or sunlight reflections. This improves segmentation of small, weak, or partially occluded fires. Attention U-Net also increases robustness under smoke and low-contrast conditions. As a result, it delivers more accurate and reliable fire masks than standard U-Net.

4.3.1 Model / Architecture

Model type: Attention U-Net (custom implementation with 4 downsampling/upsampling layers and attention gates)

Input channels: in_ch = 4 (RGB + Thermal or other 4-channel input)

Output channels: out_ch = 1 (binary segmentation mask)

DoubleConv filters: [64, 128, 256, 512]

Attention blocks: 3 Attention Gates (att3, att2, att1) applied to skip connections

Upsample mode: bilinear

Skip connections: concatenation of attention-modulated feature maps

4.3.2 Loss / Optimization

The Attention U-Net is trained using a ComboLoss, which combines Focal Loss and Dice Loss to balance class imbalance and segmentation accuracy. Focal Loss ($\alpha = 0.25$, $\gamma = 2$) emphasizes challenging fire and smoke pixels, reducing the influence of easy background pixels. Dice Loss (smooth = 1) ensures better overlap between predicted and ground-truth masks, particularly for small or fragmented fire regions. The Adam optimizer provides adaptive gradient updates for stable convergence. A learning rate of 5×10^{-4} ensures efficient training while preventing unstable parameter updates.

Loss function:

- ComboLoss = 0.5 * FocalLoss + 0.5 * DiceLoss
- FocalLoss hyperparameters: alpha = 0.25, gamma = 2
- DiceLoss smoothing: smooth = 1
- Optimizer: Adam
- Learning rate: lr = 5e-4

4.3.3 Training

Number of epochs: num_epochs = 100

Early stopping patience: patience = 15 (stop if no improvement in Dice for 15 epochs)

Device: GPU if available (cuda), else CPU

Batch size: dependent on train_loader

Binary threshold: 0.5 (sigmoid output)

4.3.4 Metrics

Dice, IoU, Accuracy, Precision, Recall, F1-score.

The performance of Attention U-Net is evaluated using multiple complementary metrics to capture segmentation quality. Dice coefficient and Intersection over Union (IoU) measure the overlap between predicted and ground-truth fire masks. Accuracy quantifies overall pixel-wise correctness, while Precision and Recall assess false positives and false negatives, which are critical for fire detection. The F1-score balances precision and recall, providing a robust indicator of mask reliability. Together, these metrics offer a comprehensive assessment of segmentation performance in challenging wildfire scenarios.

4.3.5 Workflow

1. Data Preparation

- Dataset Setup
- Downloaded RGB-T wildfire dataset (ZIP files).
- Mounted Google Drive and extracted the data.
- Dataset structure: rgb/, nir/, mask/.
- Library Installation
  - Deep learning: PyTorch, TensorFlow, Keras
  - Image processing: OpenCV, scikit-image, imageio
  - Data & visualization: numpy, pandas, matplotlib, seaborn, plotly, tqdm
- Augmentation: Albumentations
- Data Loading & Verification
- Loaded RGB, NIR, and mask images.
- Skipped missing/incomplete samples.

- Total valid samples: 1367
- Image shapes: RGB (420,420,3), (512,640,3); NIR & masks (512,640)
- Mask unique values: 0 (background), 255 (wildfire)

2. Data Preprocessing & Augmentation

- Augmentation Strategy (Albumentations)
- Training: horizontal/vertical flips, rotations, color jitter, Gaussian blur, resize, normalization
- Validation: resize and normalization only
- Custom Dataset Class
- Combined RGB + NIR → 4-channel input
- Applied augmentations dynamically
- Converted to PyTorch tensors [B,4,H,W]for images and [B,1,H,W] for masks
- Data Split & Loading
- Train: 80%, Validation: 20%
- Batch size: 4, shuffle enabled for trainingVisual Placeholder:
- Sample RGB/NIR/mask images before and after augmentation

3. Model Design

The proposed model is an Attention U-Net for binary image segmentation. It follows an encoder–decoder structure where the encoder extracts hierarchical features using Double Convolution blocks with max pooling, increasing the feature depth from 64 to 512 channels. The decoder restores spatial resolution using bilinear upsampling and feature fusion.

Attention gates are integrated into the skip connections to selectively emphasize relevant regions while suppressing background features. This improves segmentation accuracy, especially in complex scenes. The network accepts a 4-channel input and produces a single-channel segmentation mask using a 1×1 convolution followed by sigmoid activation.

4. Training & Optimization

The model is trained using a combined loss function consisting of Focal Loss and Dice Loss to handle class imbalance and improve overlap accuracy. Optimization is performed using the Adam optimizer with a learning rate of 5×10^{-4}.

Training runs for up to 100 epochs with early stopping applied if validation Dice score does not improve for 15 epochs. The best-performing model is saved based on the highest validation Dice score.

5. Evaluation & Visualization

- Loaded best saved model
- Computed confusion matrix and all metrics
- Plotted training & validation loss, Dice, IoU curves
- Displayed sample inputs, ground truths, and predicted masks Visual Placeholder:
- Learning curves (loss, Dice, IoU)
- Prediction comparison (input, ground truth, prediction)

4.3.6 Model Visualization

4.3.6.1 Architecture Visualize

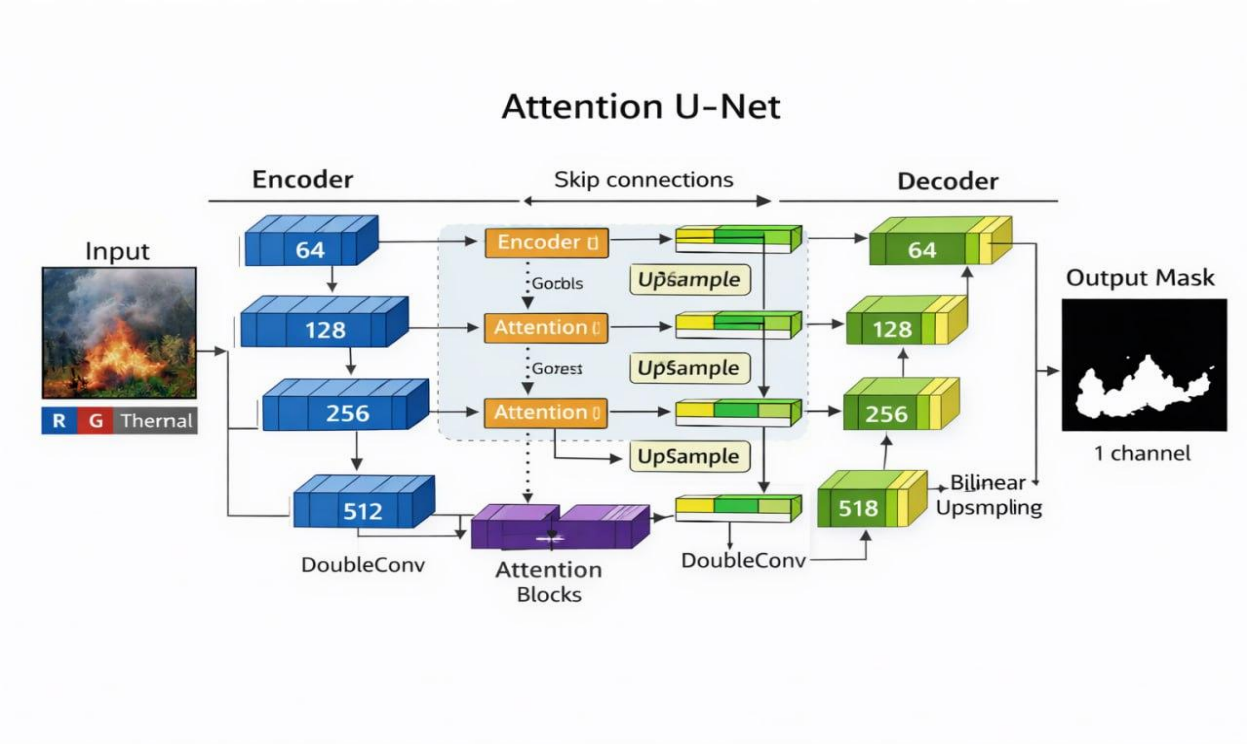

4.3.6.2 Ground Truth vs Predicted Mask

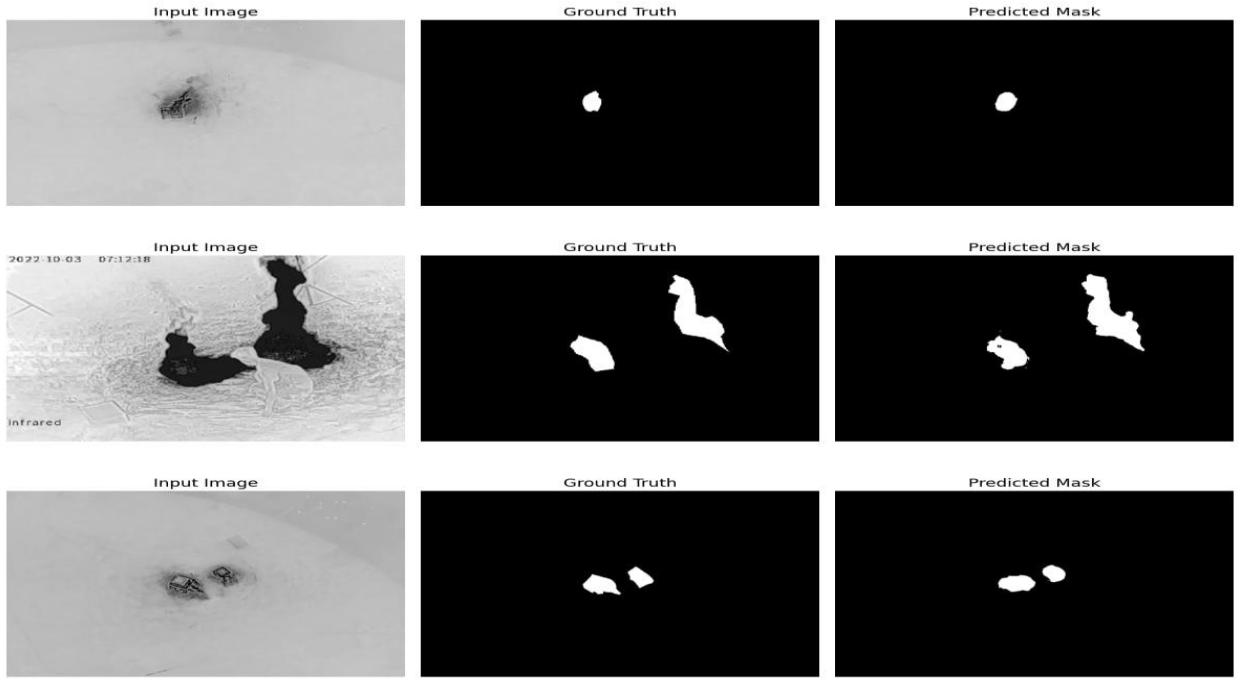

4.3.6.3 Learning Curves ,Confusion Matrix & Scores

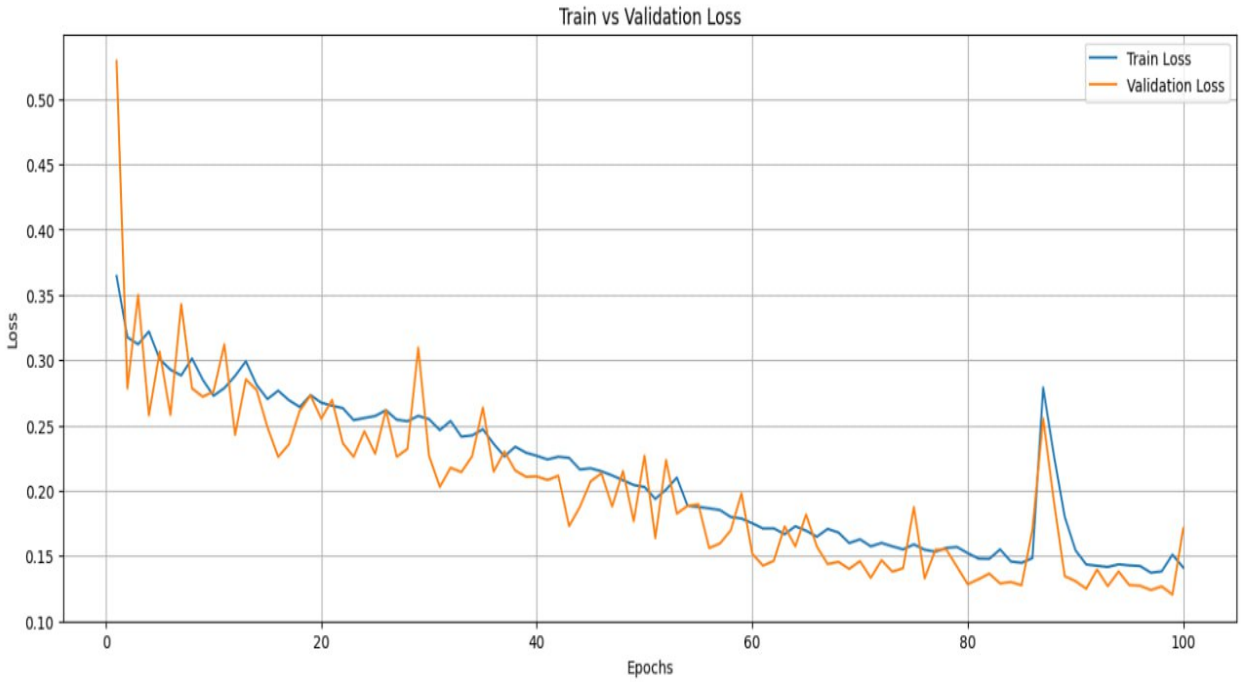

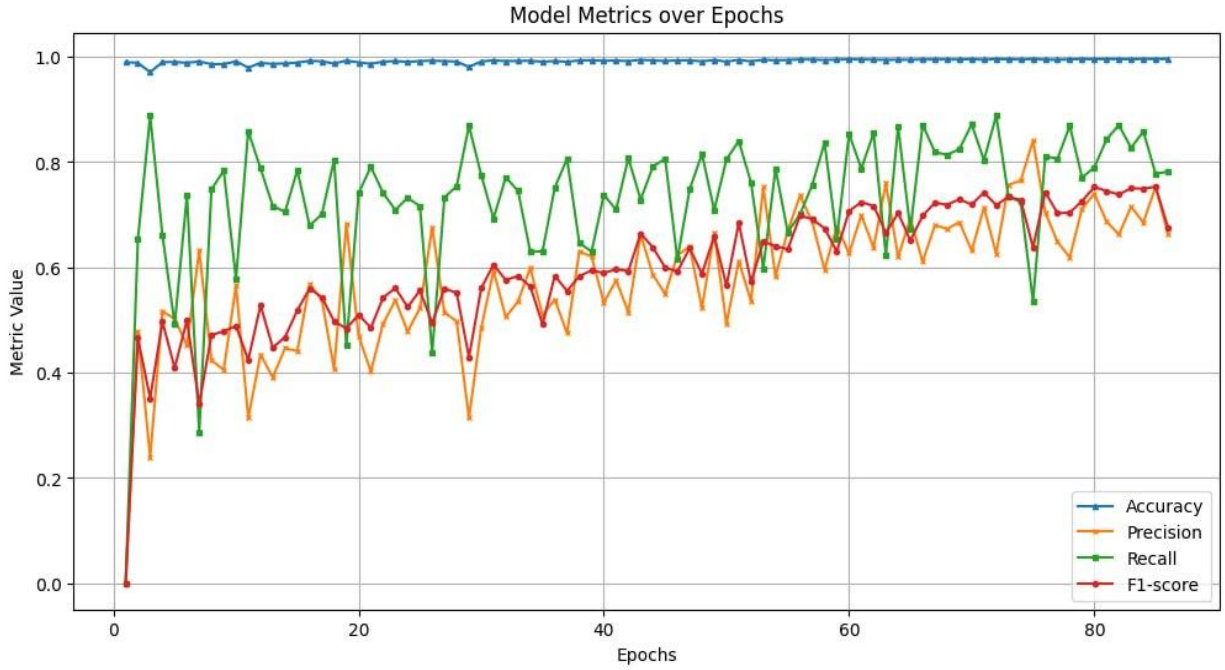

Confusion Matrix
1e7
4.0
3.5
3.0
2.5
2.0
1.5
1.0
0.5
BG
Fire
True
44381538
141228
53299
316095
BG
Fire
Predicted

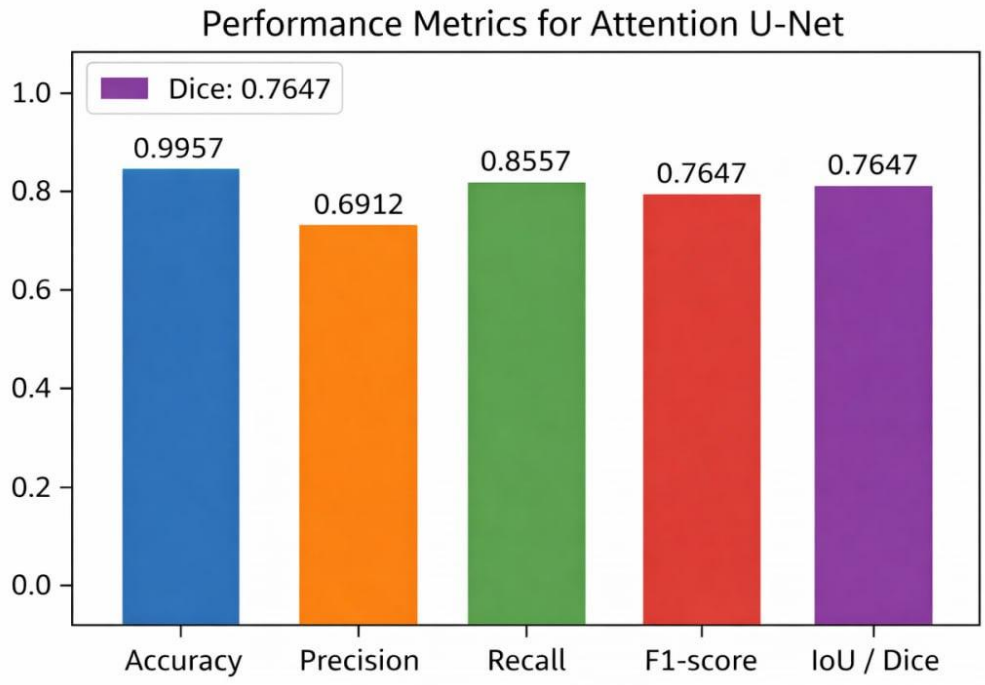
Performance Metrics for Attention U-Net
Dice: 0.7647
1.0
0.8
0.6
0.4
0.2
0.0
0.9957
0.6912
0.8557
0.7647
0.7647
Accuracy
Precision
Recall
F1-score
IoU / Dice

4.4 Models Analysis & Discussion

1. Comparative Analysis

Among the segmentation models, U-Net demonstrates superior overall performance. It achieves higher precision, F1-score, and Intersection over Union (IoU), indicating more accurate fire boundary delineation and fewer false positives. Attention U-Net slightly outperforms U-Net in recall, suggesting improved sensitivity to fire regions; however, this comes at the cost of reduced precision and overall balance.

YOLOv8 shows substantial improvement throughout training, particularly in recall and mAP@50. This indicates strong learning capability and suitability for real-time fire localization tasks. Nevertheless, its performance remains inherently less precise than segmentation-based approaches due to the complexity of object-level detection.

2. Most Important Metrics for Fire Applications

For fire segmentation, recall is the most critical metric, as failing to detect fire pixels poses a significant safety risk. IoU or Dice score follows in importance, as these metrics evaluate the spatial accuracy of fire regions. The F1-score provides a balanced measure between false alarms and missed detections.

For fire detection, recall remains the top priority, ensuring fires are not overlooked. Mean Average Precision (mAP@50) reflects overall detection robustness, while precision is essential to minimize false fire alerts.

3. Final Conclusion

Based on the presented results, U-Net is the most effective model for fire segmentation, offering the best balance between accuracy and reliability. YOLOv8 is well-suited for real-time fire detection scenarios where rapid

localization is required. For safety-critical fire monitoring systems, prioritizing recall alongside region-based metrics such as IoU or Dice is essential.

4. How to Improve Results:

• U-Net: adding BatchNorm, residual connections, and learning-rate schedulers.

• Attention U-Net: reduce attention sensitivity, adding dropout, and using pretrained encoders.

• YOLOv8: increasing image resolution, using stronger augmentation, and trying larger backbones.

4.4 Data Splitting Strategy

In this project, different data-splitting strategies were adopted based on the requirements of the tasks and the characteristics of the models used. For the wildfire segmentation task, U-Net and Attention U-Net models were trained using a three-way split consisting of training, validation, and testing sets. The validation set was used for hyperparameter tuning, monitoring convergence, and preventing overfitting through techniques such as early stopping. The test set was reserved exclusively for final performance evaluation to ensure an unbiased assessment of model generalization.

For the wildfire detection task, a YOLO-based model was trained using a train–test split. This choice was motivated by the use of fixed hyperparameters and the built-in evaluation mechanisms provided by the YOLO framework, which compute performance metrics during training. Additionally, employing a two-way split allowed a larger portion of the dataset to be utilized for training, which is particularly beneficial for object detection tasks where model performance improves with increased data availability.

Overall, the adopted splitting strategies reflect standard practices in deep learning and were selected to balance robust model evaluation with efficient use of the available data for each task.

4.5 Input Channel Configuration

Different input channel configurations were employed for the segmentation and detection models due to architectural requirements and task-specific considerations. The U-Net and Attention U-Net models were designed to process four-channel input images, where the additional channel provided supplementary information relevant to the wildfire segmentation task. Segmentation models benefit from richer input representations, as pixel-level classification requires detailed spatial and contextual cues to accurately delineate fire regions.

In contrast, the YOLO model was trained using three-channel input images (RGB). This decision was guided by the standard design of YOLO architectures, which are optimized for three-channel inputs and leverage pretrained weights learned from large-scale RGB datasets such as COCO. Introducing an additional input channel would require architectural modifications and retraining from scratch, potentially reducing training stability and eliminating the benefits of transfer learning.

Furthermore, object detection focuses on learning discriminative visual features for bounding box localization rather than precise pixel-level segmentation. As a result, the additional channel was not essential for detection performance and was excluded to maintain compatibility with the YOLO framework and ensure efficient and stable training.

4.6 Input-Level Thermal–RGB Fusion

4.6.1 Why Thermal–RGB Fusion Is Used

Wildfires exhibit strong thermal signatures that are often detectable before visible flames become prominent in RGB imagery. While RGB images capture visual characteristics such as color, texture, and smoke patterns, they are sensitive to illumination conditions, shadows, and atmospheric effects. Thermal imagery, on the other hand, directly measures heat emissions and is robust to lighting variations, enabling reliable detection of high-temperature regions even in low-visibility conditions such as nighttime or heavy smoke.

Thermal–RGB fusion combines the complementary strengths of both modalities. The RGB channels provide detailed spatial and contextual information, while the thermal channel highlights temperature anomalies associated with fire. By integrating these sources, the model gains a richer and more discriminative representation of wildfire characteristics, leading to improved robustness and accuracy, particularly for early fire detection and precise fire boundary delineation.

4.6.2 Type of Fusion Employed

The fusion approach used in this work corresponds to early (data-level) fusion. In early fusion, multiple modalities are combined at the input stage by concatenating them along the channel dimension to form a unified multi-channel representation.

In this project:

- The RGB image contributes three channels (R, G, B)
- The thermal image contributes one additional channel
- The resulting input is a four-channel (RGB + Thermal) image

Early fusion allows convolutional layers to jointly learn cross-modal feature representations from the initial stages of the network, which is particularly effective for dense prediction tasks such as semantic segmentation.

4.6.3 How Thermal–RGB Fusion Is Applied

The fusion process is implemented as follows:

1. Preprocessing and Alignment

   Thermal and RGB images are spatially aligned to ensure pixel-level correspondence. This may involve resizing, geometric correction, or registration.

2. Normalization

   Thermal values are normalized to match the dynamic range of RGB channels, ensuring stable training and balanced gradient propagation.

3. Channel Concatenation

   The normalized thermal image is concatenated with the RGB image along the channel dimension, forming a four-channel input tensor.

4. Model Input

   U-Net and Attention U-Net are modified to accept four-channel inputs by adapting the first convolutional layer.

For YOLO, which is designed for three-channel RGB input and pretrained on RGB datasets, the thermal channel is excluded to maintain architectural compatibility and leverage transfer learning.

This approach enables segmentation models to exploit thermal information at the pixel level while keeping the detection model efficient and stable.

4.6.4 Comparison with Reference Studies

Thermal–RGB fusion has been widely adopted in wildfire detection and segmentation literature. Multiple studies report that incorporating thermal data significantly improves performance compared to RGB-only approaches, particularly in challenging conditions.

Previous works have demonstrated that:

- RGB-only models may fail in scenarios with strong smoke, low contrast, or poor lighting.
- Thermal-only models lack fine spatial detail and contextual cues.
- Fused RGB–thermal models achieve superior accuracy, robustness, and generalization.

Early fusion is commonly preferred in segmentation architectures such as U-Net due to its simplicity and effectiveness in capturing spatial correspondences between modalities. Attention-based architectures further enhance this benefit by selectively emphasizing informative regions across channels, allowing the model to focus on thermally salient fire areas while preserving RGB context.

In contrast, object detection frameworks such as YOLO are predominantly designed for RGB imagery and pretrained on large RGB datasets. Several studies note that extending YOLO to additional modalities requires architectural redesign and larger datasets to avoid performance degradation, which justifies the use of RGB-only input for detection in this work.

4.6.5 Summery

Thermal–RGB fusion was employed to leverage the complementary strengths of visual and thermal information for wildfire analysis. Early fusion was used by concatenating RGB and thermal channels to form a four-channel input for segmentation models, enabling improved pixel-level fire localization. The YOLO detection model was limited to RGB input to maintain compatibility with pretrained architectures and ensure stable training. This design choice is consistent with established practices in wildfire detection research and supported by findings in prior studies.

Conclusion

This work demonstrates the effectiveness of deep learning–based segmentation and detection models for UAV wildfire monitoring using multimodal data. The use of U-Net and Attention U-Net architectures enabled accurate pixel-level fire segmentation, with the attention mechanism significantly improving focus on relevant fire regions while suppressing background noise. The YOLOv8 model further complemented the system by providing efficient real-time fire detection capabilities. Experimental results show high overall accuracy and strong Dice and IoU scores, confirming the robustness of the proposed approaches under challenging conditions. The fusion of RGB and thermal information proved crucial for early and reliable fire identification. Overall, the proposed models offer a scalable and practical solution for early wildfire detection and monitoring, supporting faster response and improved disaster management in real-world UAV applications.

References

[1]" A UAV-Based Multi-Scenario RGB-Thermal Dataset and Fusion Model for Enhanced Forest Fire Detection"
https://www.mdpi.com/2072-4292/17/15/2593

[2]" MCDet: Multi-Content Collaboration Detector for Multiscale Remote Sensing Object"
https://ieeexplore.ieee.org/document/10418876

[3]" FireMan-UAV-RGBT (IEEE Paper 10710657)"
https://ieeexplore.ieee.org/document/10710657

[4]" U3UNet — Accurate and Reliable Forest Fire Segmentation for UAV Vision"
https://ieeexplore.ieee.org/document/10944154

Appendix(1)

Mathematical equations

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

$$F1\ Score = \frac{2TP}{2TP + FP + FN}$$

$$Recall = \frac{TP}{TP + FN}$$

$$Precision = \frac{TP}{TP + FP}$$

$$IoU = \frac{TP}{TP + FP + FN}$$

$$mAP_{50} = \frac{1}{n=1} \sum_{i=1}^{n} IoU_i$$

$$Focal\ Loss = -a_{\alpha}\left(1 - exp(-\hat{y})^{T}\right)^{\gamma}\tilde{y}$$

$$Dice\ Loss = 1 - \frac{2\sum\tilde{y} + s}{\sum\tilde{y} + \sum\hat{y} + s}$$

$$Combo\ Loss = 0.5 \cdot Focal\ Loss + 0.5 \cdot Dice\ Loss$$

$$Accuracy = \frac{TP}{TP + FP + FN}$$

Appendix(2)

Terminology Table

Term	Description
UAV (Unmanned Aerial Vehicle)	An aerial platform used for remote sensing and wildfire monitoring with onboard cameras and sensors.
Wildfire Monitoring	Continuous observation of forest areas to detect, localize, and track fire or smoke outbreaks.
RGB Imaging	Visible-spectrum imaging using Red, Green, and Blue channels.
Thermal Imaging (Infrared)	Imaging modality that captures heat radiation, effective in smoke, darkness, and low visibility.
RGB–Thermal Fusion	Integration of RGB and thermal data to improve robustness and early fire detection.
Multimodal Data	Data originating from multiple sensor types (e.g., RGB + Thermal).
Image Segmentation	Pixel-level classification task that assigns each pixel to fire/smoke or background.
Binary Segmentation	Segmentation task with two classes: foreground (fire/smoke) and background.
Object Detection	Task of locating and classifying objects using bounding boxes.

Term	Description
U-Net	Encoder–decoder convolutional neural network designed for accurate image segmentation.
Attention U-Net	Modified U-Net incorporating attention gates to highlight relevant regions and suppress noise.
YOLOv8	Real-time deep learning model for object detection and segmentation on edge devices.
Encoder	Downsampling network path that extracts hierarchical feature representations.
Decoder	Upsampling network path that reconstructs spatial resolution for segmentation masks.
Skip Connections	Direct links between encoder and decoder layers to preserve spatial information.
Attention Gate	Mechanism that weights feature maps based on relevance to the target region.
Feature Maps	Intermediate representations learned by convolutional layers.
Early Fusion	Combining RGB and thermal data at the input level before feature extraction.
Feature-Level Fusion	Combining features extracted from different modalities within the network.
Focal Loss	Loss function designed to address class imbalance by focusing on hard-to-classify samples.
Dice Loss	Overlap-based loss function optimized for segmentation tasks with imbalanced classes.

Term	Description
Combo Loss	Weighted combination of Focal Loss and Dice Loss used for stable segmentation training.
BCEWithLogitsLoss	Binary cross-entropy loss combined with a sigmoid activation.
Adam Optimizer	Adaptive optimization algorithm that adjusts learning rates per parameter.
Learning Rate	Hyperparameter controlling the magnitude of weight updates during training.
Weight Decay	Regularization technique to reduce overfitting by penalizing large weights.
Batch Size	Number of samples processed in one training iteration.
Epoch	One full pass over the entire training dataset.
Early Stopping	Training strategy that halts learning when validation performance stops improving.
Accuracy	Ratio of correctly classified pixels to the total number of pixels.
Precision	Proportion of predicted fire pixels that are actually fire.
Recall	Proportion of actual fire pixels that are correctly detected.
F1-Score	Harmonic mean of precision and recall, balancing false positives and false negatives.

Term	Description
IoU (Intersection over Union)	Ratio of overlap between predicted and ground-truth regions.
Dice Coefficient	Similarity metric measuring overlap between predicted and true segmentation masks.
mAP50	Mean Average Precision at IoU threshold of 0.5, commonly used in object detection.
Ground Truth Mask	Pixel-level annotation indicating true fire or smoke regions.
Label Imbalance	Disproportion between fire and background pixels in wildfire datasets.
Real-Time Inference	Ability of a model to process data fast enough for live UAV deployment.
Edge Deployment	Running models on resource-constrained devices such as UAV onboard computers.
FireMan-UAV-RGBT Dataset	UAV-based RGB–Thermal wildfire dataset used for detection and segmentation research.
Dataset Generalization	Ability of a trained model to perform well on unseen environments.
False Positives	Background regions incorrectly classified as fire or smoke.
False Negatives	Fire regions missed by the model.
Early Wildfire Detection	Identification of fire at an initial stage before large-scale spread.